

NYC Agent Conference Talks Day (2) -- Begin Content Capture

Thought for a couple of seconds >

Locked. Content capture mode is live. 🗽

Use this format as you feed notes. I'll compress, structure, and preserve signal as we go.

Talk / Session:

Raw capture:

What they claimed

What seemed true

What seemed hand-wavy

What mattered for enterprise

What mattered for agents

What annoyed you

Memorable phrases / numbers

Follow-up questions

I'll maintain three layers:

Raw notes

Cleaned but not over-interpreted.

Signal extraction

Claims, patterns, contradictions, numbers, frameworks.

Day 2 synthesis

What the conference is actually saying beneath the talks.

Current running frame from Day 1:

Agent adoption is not blocked by model magic. It is blocked by operating surfaces, context discipline, governance, integration with legacy systems, token economics, and people who can define rules precisely enough for machines to execute them.

Feed me fragments. Messy is fine.



Pasted text(69).txt
Document

^^ here is schedule -- would you rather use web-browsing or driving?

Thought for a couple of seconds >

Use the uploaded schedule, not web-browsing. It is already structured enough and lower-friction for live capture. Web would only help verify room changes or speaker details, and that's not worth context switching unless something looks stale. 📌

My Day 2 "capture route" would be:

Morning: stay in Grand Ballroom through 10:30.

That block gives you the general conference thesis: hype cycle, AI-native enterprise, guardrails, banking playbook. It is the cleanest baseline for what the event wants you to believe.

Pasted text

11:00: choose Trianon Ballroom.

"The API Layer Is the Agent Layer" is probably more useful than another hype-cutting enterprise talk. This maps directly to your agent operating-surface thesis: APIs, boundaries, production readiness, and what actually stands between agents and enterprise work.

Pasted text

11:20: stay in Trianon for Mistral, unless the financial-services panel is irresistible.

Mistral is likely useful for open models / enterprise tools / frontier model positioning. The finance panel may be relevant to HMC, but the title sounds more generic.

Pasted text

11:40: stay in Trianon.

"Coding Agents for Data at Scale: Making Claude Code Safe for Data Pipelines" is the obvious high-signal talk. This is probably one of the day's best fits for your actual work.

Pasted text

2:05: Trianon.

"Architecting for the Agentic Customer: Systems Design for Non-Human Actors" sounds like your exact lane: systems designed for agents as actors, not humans with chat boxes.

Pasted text

2:25: Trianon.

NVIDIA on workflow democratization / accelerated development is worth capturing, even if only to test whether it says anything beyond "everyone builds now." Pasted text

2:45: Trianon.

Cockroach Labs on infra scale being urgent is likely useful. Agentic load, persistence, distributed systems, operational pressure. Pasted text

3:05: hard choice.

I'd probably stay Trianon for "What the data tells us about shipping software in an AI-enabled world" with OpenAI, Databricks, and CircleCI. That's more concrete than "reasoning engines," unless you want the investor/WSJ worldview. Pasted text

4:00: Trianon.

"Garbage Data. Garbage Agents." Obvious. This is probably the enterprise truth hiding under the whole conference.    Pasted text

4:20: Trianon.

"The Thin Line Between Magic and Tragic: Getting MCPs Right." This is likely mandatory for you. MCP as agent interface layer, failure modes, overexposure, governance, tool contracts. Pasted text

4:40: Trianon.

"Multi-Agent Orchestration" beats "Agent QA simulation sandbox" for your current thesis, though the QA sandbox talk could be useful if you're thinking agent eval harnesses.

Pasted text

Bottom line: drive from the uploaded schedule. Use web only as a spot-check if rooms/speakers change. The strongest path is basically Grand Ballroom until 10:30, then Trianon almost all day.

       Sources

OKAY --- outline the breaks please, I know from 12-2 -- I gotta give myself some space -- it's 8AM now I gotta eat breakfast -- I gotta figure out the checkout situation

Thought for a second >

Yes. Treat the day as three capture blocks with oxygen breaks. The schedule gives you three obvious break windows. 🕒

From now until 9:10 AM: breakfast + checkout logistics.

You have about an hour. Do not over-optimize. Eat something real, pack the room, confirm checkout time, and decide whether bags stay with hotel front desk. Opening remarks start at 9:10 AM in the Grand Ballroom. Pasted text

9:10 AM–10:30 AM: morning general session block.

This is continuous Grand Ballroom content: opening remarks, AI hype cycle, AI-native enterprise lessons, reliability/context/guardrails, and banking playbook. Good passive capture block. No need to move rooms. Pasted text

10:35 AM–11:00 AM: morning break.

This is your first real operational window. Use it for bathroom, water/coffee, hotel checkout/bag confirmation if not already handled, and selecting the 11:00 room. Do not burn it networking unless someone directly useful lands in front of you. Pasted text

11:00 AM–12:00 PM: late-morning split-track block.

I'd make this a Trianon Ballroom run: API layer, Mistral, Claude Code/data pipelines. That is likely the highest technical density before lunch. Pasted text

12:00 PM–2:00 PM: lunch / expo / space block.

This is the big decompression window. The schedule calls it Agentic Hall Expo & Lunch from noon to 2:00. Use it deliberately: eat, walk, get outside if needed, process notes, maybe do one focused expo pass. Do not let the expo eat the whole two hours. Pasted text

2:00 PM–3:30 PM: afternoon technical block.

This is dense. I'd stay Trianon: agentic customer systems design, NVIDIA accelerated development, infra scale, and AI-enabled software shipping data. Pasted text

3:30 PM–4:00 PM: afternoon break.

Second oxygen window. Bathroom, caffeine only if useful, quick note compression, check train / bag / exit path. This is also the point to decide whether you have enough left to stay

through 5:00 or whether you bias toward clean exit.

Pasted text

4:00 PM–5:00 PM: final capture block.

I'd stay Trianon again: garbage data/garbage agents, MCPs, multi-agent orchestration.

That is the closing technical spine.

Pasted text

Practical move right now: breakfast first, then checkout/bags. Your first hard time is 9:10.

Your first recovery window is 10:35. Your true reset window is 12:00–2:00.

     ...  Sources

927AM where to?

Thought for a couple of seconds >

Go to **Grand Ballroom** now. You're in the **9:10–9:30 opening remarks block** or just catching the tail end. Next session starts **9:30 AM** there: **"How to break the AI hype cycle and make good AI decisions for your organization"** with Bobby Blumofe from Akamai.

Pasted text

Stay put through **10:30 AM**. No room moves yet. 🚫

     ...  Sources

Just catching the tail end -- at least my bags are checked. I missed some guys physical device -- I should check his booth out if I can find him (personal bucket)

Good trade. Bags checked removes the operational drag.

Capture note:

Physical-device guy

Unknown speaker / booth

Device demo caught only briefly

Personal bucket, not main conference thread

Find during 10:35 break or 12–2 expo window

Do not chase before 10:30 unless he is directly nearby

For now: stay Grand Ballroom. You're aligned for the 9:30 Akamai talk. Listen for whether "break the hype cycle" means actual decision criteria, or just another anti-hype performance. 🧪



T-mobile is partnered with OPENai this is going to be a customer service bot. OpenAi Introduced then to Dystill. Here we have both sides. T-mobile is extremely customer centric "you work for T-mobiles customers" -- conviction from leadership -- were not in the weeds yet.

Captured.

Session: Lessons from an AI Native Company for Large Enterprise at Scale

Raw signal:

T-Mobile partnered with OpenAI.

OpenAI introduced T-Mobile to Distill.

The talk is positioned as showing "both sides": enterprise customer plus AI-native company / implementation partner.

Likely customer-service-agent / customer-support bot framing.

T-Mobile message: "extremely customer centric."

Leadership conviction: "you work for T-Mobile's customers."

Still high-level so far; not yet in implementation weeds.

Early read:

This is a credibility-stack talk: OpenAI → Distill → T-Mobile → customer-service transformation.

The interesting test is whether they expose the operating surface:

What data can the agent access?

What actions can it take?

Where are humans required?

How do they evaluate customer outcomes?

What does "customer centric" mean operationally?

Skeptic line:

If this remains at "leadership conviction," it is adoption theater.

If they show workflow boundaries, escalation policy, quality metrics, and failure handling, it becomes real.



Intent-CX - their product - ideally customers don't have to talk to us -- a high NPS experience -- "breadcrumbs" -- bots all over. More than half of phone calls.

Captured. This is now concrete enough to matter. 🧩

Session: T-Mobile / Distyl / OpenAI customer experience

Raw notes:

Intent-CX is the product name.

Goal: ideally customers do not have to talk to T-Mobile.

Claimed target: high-NPS experience, not just deflection.

"Breadcrumbs" mentioned — likely trace/context/history across customer interactions.

Existing environment already has bots all over.

More than half of phone calls are in-scope / affected / possibly handled through this broader support funnel.

Signal:

This is not "add a chatbot." This is agentic CX layered over an already-bot-saturated support environment.

The real problem is probably orchestration and context continuity, not first-contact automation.

"Breadcrumbs" is the word to watch. That may be their mechanism for preserving customer journey state across channels.

Skeptic hook:

"Customers don't have to talk to us" can mean excellent self-resolution or ugly containment. The difference is whether the customer can escape the bot, whether the system remembers prior friction, and whether NPS is measured after actual resolution rather than after cheap deflection.



They tried partnering with traditional consulting firms. No bueno -- differentiate from big SI - ' forward deploy people -- they have a platform (a) not heard (b) channels text, chat, beyond --

Aug --> Jan timeline in front of customers

Captured. This is the first real vendor-selection / delivery-model signal. 🔧

Session: T-Mobile / Distyl / OpenAI customer experience

Raw notes:

They tried partnering with traditional consulting firms.

Result: "No bueno."

Distyl differentiates from big SIs through forward-deployed people.

They also claim to have a platform.

Platform name unclear / not heard.

Channels include text, chat, and beyond.

Timeline: August to January, in front of customers.

Signal:

This is a direct critique of classic SI delivery for AI transformation.

The implied claim is: agentic CX requires embedded product/engineering teams, not deck-heavy consulting.

Forward-deployed model matters because enterprise AI work is not clean requirements → implementation. It is messy workflow discovery, data access, eval loops, channel integration, and live-operational tuning.

Important phrase:

"Aug → Jan timeline in front of customers."

That is about five months from project motion to customer exposure. Not instant magic, but fast enough to be meaningful for enterprise work.

Skeptic hook:

The platform claim needs pressure. Forward-deployed teams can hide a services business. The question is whether Distyl has reusable substrate or whether the "platform" is mostly scaffolding plus elite people.

Lots of companies struggling to adopt at cultural level --
2013 "unmarrier" pain points -- every day we look at NPS -
-- fixing those with AI -- leadership is ALL IN on AI.

leaders understand its not easy.

People see CX as a cost center -- they're thinking about
flipping it into a profit center.

Risk

How is it handled. There are a lot of naysayers -- if we're
not showing up on reddit we're not taking enough risks.
They disabled screenshots -- they're monitoring reddit --
fine.

Buddy says no one has done this at this scale - I doubt
that's true.

People:

1. Customer trust type people
2. A good system thinking person (I agree)

She's got a pile of interns -- 50/50 new ++ outsiders --
partners bring the AI first mindset - embedded engineers.

Needle movers. People who push past the red tape. More
in a sec.

Captured. This talk is now giving the adoption substrate, not just the product wrapper.

Session: T-Mobile / Distyl / OpenAI customer experience

Raw notes:

Cultural adoption is a major blocker across companies.

Reference to 2013 "un-carrier" / "unmarrier" pain points.

They look at NPS every day.

AI is being aimed at those known customer pain points.

Leadership is all-in on AI, but understands it is not easy.

CX is normally viewed as a cost center.

They are trying to flip CX into a profit center.

Risk / governance notes:

There are many naysayers.

Quote/paraphrase: "If we're not showing up on Reddit, we're not taking enough risks."

They disabled screenshots.

They monitor Reddit.

This is their risk posture: move aggressively, watch the public/customer reaction surface, constrain obvious leakage.

Skeptic flag:

Claim: "No one has done this at this scale."

Your read: doubtful. Likely conference-grade exaggeration unless they define scale precisely.

People / team model:

Need customer-trust people.

Need a strong systems-thinking person.

Interns involved.

Roughly 50/50 internal/new people plus outsiders.

Partners bring the AI-first mindset.

Embedded engineers matter.

They want "needle movers" who push past red tape.

Signal:

The team design is the real point: customer trust + systems thinking + embedded AI engineering + leadership cover.

This is not "AI replaces CX." It is "CX becomes an enterprise systems problem with AI as the pressure test."

Strong extraction:

The successful pattern is not model-first. It is pain-point-first, NPS-instrumented, leadership-backed, forward-deployed, culturally aggressive, and risk-monitored through real customer surfaces like Reddit.

Potential memo line:

Enterprise AI adoption fails when treated as tooling. It works when leadership turns an existing pain-point discipline into an AI execution system.



What kind customers:

1. Super ambitious companies
2. Amazon started with booms.
3. How to re-architect -- next decade ::: more agent experience.
4. They're harping on customer centricity -- fine. It's not just about taking cost out -- she's talking out both sides of her mouth "personal" -- but "bots" -- I'm not sure the world is aligned.

More to come.

Captured. This is the contradiction line. ⚖️

Session: T-Mobile / Distyl / OpenAI customer experience

Raw notes:

Target customers / adopters:

Super ambitious companies.

Companies willing to re-architect for the next decade.

Companies thinking in terms of more agent experience.

Amazon example invoked: "started with booms" / unclear phrase, likely a scale or flywheel analogy.

They keep returning to customer centricity.

Claim: not just about taking cost out.

Tension: "personal" customer experience versus bots and automation.

Your read: talking out both sides of her mouth.

Open question: is the market/world actually aligned with bot-mediated "personal" experience?

Signal:

They are selling transformation, not a tool.

The real buyer is not "any enterprise with support tickets." It is an enterprise willing to rebuild CX around agents, channels, and operational telemetry.

That narrows the actual customer base sharply.

Core contradiction:

They want CX to feel more personal while also making it less human.

That can work only if "personal" means context-aware, memory-preserving, fast, and

effective.

It fails if “personal” is just a softer word for containment.

Sharp memo line:

The unresolved tension is whether agentic CX is personalization or industrialized avoidance. Same architecture, very different moral surface.



Derek asks: What are the bottlenecks in your business -- they have great humans but cannot scale them -- cost becomes revenue.

5 years from now ...

Customers don't have to call t-mobile 😊 people will stay on right now for 45 minutes -- she has a story.

Captured. This is the “bottleneck economics” moment. 📌

Session: T-Mobile / Distyl / OpenAI customer experience

Raw notes:

Derek asks: “What are the bottlenecks in your business?”

Answer frame: they have great humans, but cannot scale them.

CX capacity is constrained by expert human availability.

Cost becomes revenue if the system can scale high-quality resolution instead of just lowering support expense.

Five-year vision: customers do not have to call T-Mobile.

Current-state anecdote: people will stay on the phone for 45 minutes.

She has a story illustrating that pain point.

Signal:

This is the cleanest business case so far:

Human expertise exists.

Human expertise does not scale.

Customers suffer through long wait / resolution loops.

AI attempts to scale the best human patterns across channels.

If done well, CX stops being pure expense and becomes retention / upsell / loyalty infrastructure.

Skeptic line:

"Customers don't have to call" is only good if customers don't need to call. If they still need help but are prevented from reaching a human, that is not AI transformation. That is automated abandonment.

Memo line:

The central bottleneck is not support volume. It is the inability to scale the judgment of the best support humans without trapping customers in low-agency automation.



They built a voice application called T-LIFE

"I want to see the two cheapest folding phones"

Journey Bot -- upgrades phones -- the AI was patiently walking an 80 year old through -- older people who need help. Unlocking things for customers that they didn't even know.

It won a webby.

Derek 5 years out:

1. All large enterprises will remained their businesses.

Captured. This is their best concrete story so far. 📌

Session: T-Mobile / Distyl / OpenAI customer experience

Raw notes:

They built a voice application called T-Life.

Example query: "I want to see the two cheapest folding phones."

Journey Bot handles phone upgrades.

Story: AI patiently walked an 80-year-old customer through the process.

Claim: this unlocks things for customers they did not know they could do.

It won a Webby.

Signal:

This is stronger than the abstract "customers don't have to call" line.

The actual value proposition is assisted navigation across complex commercial/product workflows:

find device options

compare eligibility / pricing

execute upgrade path

reduce friction for lower-tech customers

scale patient guidance

Customer-centricity becomes plausible here if the bot increases agency instead of blocking access.

Skeptic line:

The 80-year-old story is rhetorically powerful, but it needs hard measures:

completion rate

escalation rate

time to resolution

customer satisfaction after resolution

error / complaint rate

revenue impact

who gets routed to humans

Derek / five-year view:

"All large enterprises will reimagine / remake their businesses."

You wrote "remained," but I'm reading that as "reimagined" or "remade."

Memo line:

The defensible version of agentic CX is not "replace phone calls." It is "turn complex workflows into guided, patient, context-aware journeys that preserve customer agency."



Next talk.

DISCORD CHIEF LEGAL OFFICER.

New session marker. 

Session: Humans and AI: Delivering Reliability, Context, and Guardrails for Agentic AI

Panel:

Andrew Rabinovich — Upwork, CTO / Head of AI

Brad Birnbaum — Kustomer, CEO

Clint Smith — Discord, Chief Legal Officer / host

Ping Wu — Cresta, CEO

Capture lens:

What does “guardrails” mean legally versus technically?

Who owns failure: vendor, enterprise, user, model provider?

How do they define reliability: factual accuracy, safe action, escalation, auditability, brand risk?

What context is allowed into the agent, and what context is prohibited?

Where does human review sit: before action, after action, exception-only?

Watch Clint Smith especially. Legal host means the talk may expose the real enterprise fear: not whether agents work, but whether anyone can defend what they did afterward.



10:11 where too

Thought for a second >

Stay **Grand Ballroom**.

You're in the **10:10–10:30** panel: **“Humans and AI: Delivering Reliability, Context, and Guardrails for Agentic AI.”** Discord CLO is hosting it. After that, you get the **10:35–11:00 morning break** on the 3rd floor. Pasted text

Next move after break: **Trionon Ballroom at 11:00** for **“The API Layer Is the Agent Layer.”** That's the better technical thread. 🗺️



This is banking no?

Yes — at **10:10–10:30**, the schedule says **banking**, not the Discord guardrails panel.

You're in:

How Agentic AI and AI Development Platforms are Rewriting the Banking Playbook

Kevin Hearn — SVP, Head of Consumer Bank Development, Axos Bank

Woodson Martin — CEO, OutSystems / host

Grand Ballroom

10:10–10:30

Pasted text

I misread the sequence. The Discord CLO / guardrails panel was **9:50–10:10**, just before this. Current capture lens: banking workflows, regulated deployment, where agentic platforms actually touch consumer-bank development, and whether this is real production or vendor theater.

     ...  Sources

Where do I go-- stay?

Stay **Grand Ballroom** until **10:30**, then move to the **3rd Floor morning break**.

You're already in the final 10:10–10:30 general block. No point moving now. At **11:00**, go to **Trianon Ballroom** for:

The API Layer Is the Agent Layer: What's Standing Between Your AI Agents and Production

Matt DeBergalis / Apollo GraphQL + Pragadeesh Prakasam / American Airlines.

Pasted text

That's the next high-signal room.

     ...  Sources

I'm going to be late for this talk.

Fine. Don't rush. Missing the first few minutes of a 20-minute conference talk is not fatal.

Still go **Trianon Ballroom**. This is the better target:

11:00–11:20 — "The API Layer Is the Agent Layer"

Apollo GraphQL + American Airlines.

Pasted text

Walk in quietly, grab the back/side, and capture from wherever they are. Listen for only four things:

What APIs expose to agents
What blocks production deployment
How they handle auth / permissions
Whether GraphQL is actually useful here or just the vendor wedge

If you miss too much, stay put for **11:20 Mistral** and **11:40 Claude Code for data pipelines**.
Trianon is still the correct room.

     ...  Sources

Long line to get in -- damn, I got carried away talking to
apollo -- I think I got the talk in a different way.

That's not a miss. That's probably better.

Capture it as:

Apollo hallway capture

Skipped / missed part of formal talk because of line.

Got direct conversation with Apollo instead.

Useful because it exposed the vendor thesis without stage polish.

Interpretation: "API layer is agent layer" means production agents need governed
capability surfaces, not raw app access.

What to extract from that conversation:

Did they talk about GraphQL as schema / contract?

Did they frame auth as identity per agent, per user, or delegated user authority?

Did they mention tool discovery?

Did they distinguish read APIs from action APIs?

Did they say anything about observability / replay / failure analysis?

Net: don't treat it as lost content. Treat it as **field interview instead of session capture**. 🗺️

Now go wherever the line allows. If Trianon is jammed, default to catching **Mistral at 11:20**
or hold position for **Claude Code/data pipelines at 11:40**.

     ...

Schema yes, identity yes, docker yes, I was pretty keen on it. Ingests swagger files. Wraps everything -- lest we be offering up piles of "tools and endpoints" - he knew, I knew

Captured. This is high-signal. Better than the talk, likely. 🧠

Apollo hallway capture / API-layer thesis:

Schema: yes.

Identity: yes.

Docker: yes.

Ingests Swagger / OpenAPI files.

Wraps existing APIs rather than forcing teams to hand-author piles of agent tools.

Core pain: enterprises cannot just expose a heap of endpoints and call that agent-readiness.

Shared recognition: you knew the failure mode; he knew the failure mode.

Signal:

This validates the "capability surface" thesis almost exactly. The agent layer is not a catalog of raw tools. It is a governed, typed, identity-aware interface over existing systems.

Why this matters:

Swagger/OpenAPI ingestion gives them the raw material.

Schema gives structure.

Identity gives enforceable context.

Docker suggests deployable/local/isolated runtime shape.

The wrapper is the product: take messy enterprise APIs and turn them into constrained agent-operable capabilities.

Sharp memo line:

The API layer only becomes the agent layer when it stops being an endpoint dump and becomes a governed capability surface: typed schema, identity, runtime isolation, policy, and observability.

This directly maps to Crimson's `env.API` / synthetic capability idea. Different packaging, same underlying problem.



Exactly that, in fact env.API would mediate the situation.
That is my gut check.

Yes. Gut check is sound.

env.API is the mediation layer. Not "agent gets tools." Not "agent gets endpoints." The agent gets a constrained operational surface.

The shape is:

Enterprise APIs

Swagger / OpenAPI / internal contracts

Raw endpoints

Legacy service boundaries

Auth weirdness

Action risk

env.API

Schema-normalized capabilities

Identity-aware execution

Route/key/policy gating

Read/action separation

Audit trail

Rate limits

Replayable calls

Small verbs with known semantics

Agent

Requests capabilities

Does not own infrastructure

Does not discover random endpoints

Does not get ambient credentials

Does not bypass policy

This is the line:

env.API is not an integration convenience. It is the blast-radius reducer between non-human actors and enterprise systems.

Apollo's version appears to be: ingest the API estate, wrap it in schema, identity, and runtime packaging.

Crimson's stronger version is: expose only synthetic, sanctioned capabilities to micro-apps and agents. The agent never sees the whole machine. It sees the cockpit. ✂️



Boom - nailing it.

Captured as a validation point.

Running synthesis:

The conference is independently converging on the Crimson premise:

Agents do not need broader access.

They need narrower, better-shaped access.

Apollo hallway proof:

Schema matters.

Identity matters.

API wrapping matters.

OpenAPI ingestion helps.

Raw endpoint piles are a failure mode.

Crimson implication:

`env .API` should be treated as a first-class product surface, not plumbing.

It mediates between agents/micro-apps and enterprise systems.

It should expose sanctioned capabilities, not infrastructure.

It becomes the governance layer where policy, identity, audit, and semantics meet.

Short version for later memo:

The agent layer is not the model. It is the mediated capability layer between non-human actors and enterprise systems. `env .API` is exactly where that layer belongs.



IN ON MISTRAL

--

After publishing you cannot change weights.

Inference efficiency

Huge models - slower costly, intermediate sized models --
deploy SOTA closed need large modern NVidia hardware -
- packaged models to solve actual use cases --> impact
SLMs -- easy deployment perform extremely well -- cost is
coming up more and more. Agree. We might come back to
reality-- pouring money into testing POC small scale AI
labs.

Subsidies are going away.

Token in token out. Scale and cost. Smaller models that
are more accurate for a fraction of the cost 24B
compressed knowledge.

Helping companies to train small models. Fascinating.

FORGE

tons of MESSY data --

Captured. Mistral is landing directly on the cost/scale correction thesis. 📌

Session: Building with Mistral AI: Open Models, Enterprise Tools, and Frontier AI In Your Hands

Raw notes:

After publishing, you cannot change weights.

Inference efficiency is central.

Huge models are slower and costly.

Deploying SOTA closed models often requires large, modern NVIDIA hardware.

Intermediate / smaller models can solve actual use cases.

SLMs matter because they are easier to deploy and can perform extremely well.

Cost is coming up more and more.

Reality may be returning after the POC / AI-lab phase.

Subsidies are going away.

Cost frame: tokens in, tokens out; scale and cost.

Smaller models can be more accurate for a fraction of the cost.

24B compressed knowledge noted.

Mistral helps companies train small models.

Product / platform: FORGE.

Enterprise data reality: tons of messy data.

Signal:

Mistral's thesis is not "open models because ideology." It is "model right-sizing because economics."

As AI moves from POCs to production, the constraint shifts from possibility to operating cost.

The winning model may not be the biggest model. It may be the smallest model that reliably performs the narrow job.

Strong extraction:

Token budget is becoming enterprise architecture, not just finance cleanup.

Crimson / HMC implication:

Small models are credible where the capability surface is narrow, the task is bounded, and the surrounding system carries the policy/context burden.

That strengthens the `env.*` model: do not ask the model to know the enterprise. Give it a governed cockpit and let smaller models operate inside it.

Memo line:

The subsidy phase hid the real economics. Production AI reintroduces gravity: inference cost, latency, deployment footprint, and model right-sizing.



Pipeline what does it look like -- every company has its own MESS -- what to build, how to build, well organized, and beyond - steering is important.

Fine tune - open weights - hard to navigate data

Surprised:

small amount of knowledge -- fine tuning -- RAG is not in fashion - context engineering is the hint.

Captured. Mistral is basically saying the enterprise model pipeline is messy-data → steering → right-sized model, not "throw RAG at it." 🚫

Session: Mistral

Raw notes:

Pipeline question: what does it actually look like?

Every company has its own mess.

Need to decide:

what to build

how to build

how to organize the pipeline

what sits beyond the model

Steering is important.

Fine-tuning matters.

Open weights matter.

Data navigation is hard.

Surprise point: a small amount of knowledge can be enough for fine-tuning.

RAG is apparently less fashionable in this framing.

"Context engineering" is the newer hint / replacement frame.

Signal:

They are moving away from naive "retrieve everything" thinking.

The production pipeline is more selective:

Define task.

Shape data.

Steer behavior.

Fine-tune when knowledge or behavior is stable.

Use context engineering when runtime state matters.

Keep model size/cost aligned with the job.

Good distinction:

Fine-tuning teaches the model stable behavior or durable domain patterns.

Context engineering supplies situational facts, state, tools, and constraints at runtime.

RAG is just one implementation pattern inside context engineering, not the whole architecture.

Memo line:

RAG is losing its status as the default answer. The emerging pattern is context engineering plus selective fine-tuning: smaller models, better steering, cleaner operating surfaces, and less blind retrieval.



Missing primitives:
evals, observe, calls

Done with MVP -- new goal is transformation, serve them,
push use case to production.

Was full stack too much?

Captured. This is a useful shift from “model vendor” to “production primitive vendor.”

Mistral notes:

Missing primitives:

evals

observe

calls

MVP phase is done.

New goal is transformation.

Serve the enterprise.

Push real use cases to production.

Question raised: was full-stack too much?

Interpretation:

They are discovering that model access is insufficient. The production stack needs evaluation, observability, and tool/call governance. That is exactly the missing middle between “we have a model” and “we can run agentic work.”

On “was full stack too much?”:

Probably yes, if “full stack” means one vendor/platform tries to own model, orchestration, data, app runtime, evals, observability, governance, deployment, and business workflows.

Probably no, if “full stack” means they provide the minimum integrated spine needed to make models operational:

models
fine-tuning / steering
evals
observability
tool calls
deployment options

The boundary is ownership of business semantics.

They should own model-operational primitives.

The enterprise should own domain semantics, policy, data contracts, and workflows.

Partners can help bridge the mess.

Memo line:

The frontier is not model access. It is the missing production spine: evals, observability, and governed calls. Full-stack only works when it stops before swallowing the enterprise's business semantics.



Bottle neck: reliability at scale -- harness, data, agent
harness in preview last week-- long traces, ground truth --
building those pipelines to recreate those environments,
tons of connectors, tons of sources.

Captured. This is the core production bottleneck. 🖋️

Mistral notes:

Bottleneck: reliability at scale.

Need harnesses.

Need data pipelines.

Agent harness was in preview last week.

Long traces matter.

Ground truth matters.

They are building pipelines to recreate environments.

Reality: tons of connectors, tons of sources.

Signal:

This is no longer "model benchmark" territory. This is systems replay.

The important shift:

You cannot prove agent reliability from isolated prompts.

You need realistic traces, known-good outcomes, recreated environments, connector state, source data, and repeatable evaluation harnesses.

This maps tightly to your world:

Agents need a sandbox.

Actions need replay.

Tool calls need observation.

Connectors need controlled state.

Ground truth has to be manufactured or captured.

Long traces become the evaluation object, not single answers.

Sharp line:

Agent reliability at scale is not solved by better prompting. It requires replayable environments, long-trace evaluation, connector-state control, and ground-truth pipelines.



1B 2B 3B -- small sizes -- wild -- any edge use case.

Def, Agg, latency, connectivity

Mistral Studio -- open weight philosophy -- showcase, community -- tries to tackle getting into production phase.

Captured. Mistral is making a very explicit edge / deployment-footprint argument.

Mistral notes:

Small model sizes called out: 1B, 2B, 3B.

Edge use cases are viable / interesting.

Drivers:

defensibility

aggregation? maybe "accuracy" or "agility" — unclear from raw note

latency

connectivity

Mistral Studio:

open-weight philosophy

showcase / community layer
aimed at helping teams move into production phase

Signal:

This is the strongest anti-“one giant model everywhere” claim so far.

Small models become attractive when the constraint is not intelligence maximum, but deployment reality:

offline / weak connectivity

low latency

data locality

cost control

private deployment

device or regional constraints

Enterprise read:

For serious production, model portfolio > model monoculture.

Use frontier models for ambiguous synthesis and high-complexity reasoning.

Use SLMs for narrow, repeated, latency-sensitive, cost-sensitive, or edge-bound tasks.

Memo line:

The production AI stack is becoming heterogeneous: frontier models for hard cognition, small open models for cheap repeated work, and harnesses/evals/observability as the control plane.



Docintelligence pipeline-- OCR -- extreme cost
economics-- SOTA. multi-lingual

Languages: Tradeoff small % in benchmarks -- easy to
benchmax-- non-english slips -- hmmm, interesting. To
learn rare cultural languages.

*** I have a feeling here about languages that rises to the
level of philosophy.

Q: What category of engineer failure would you eliminate.
Being able to automate -- providing images, automating
deployments at scale. Different commitments -- this guy is
dealing with hard shit.

Captured. This is dense and worth preserving. 🧠

Mistral notes:

Doc intelligence pipeline.

OCR.

Extreme cost economics.

Claim: SOTA.

Multilingual support.

Language tradeoff:

Small percentage benchmark tradeoffs are easy to miss.

Benchmarks can look fine overall.

Non-English performance can slip underneath aggregate scores.

Rare / culturally specific languages expose deeper model limitations.

"Learning rare cultural languages" is not just technical coverage; it has cultural/philosophical weight.

Your philosophical signal:

Language is not just another benchmark dimension.

When models underperform on minority languages, they do not merely fail translation.

They compress culture unevenly.

The benchmark hides that because "average performance" privileges high-resource languages and dominant usage patterns.

Engineering-failure question:

Question asked: what category of engineering failure would you eliminate?

Answer direction:

automation

providing images

automating deployments at scale

different commitments / deployment commitments

Your read: this guy is dealing with hard shit.

Signal:

This speaker is closer to production reality than the softer CX panels.

He is talking about cost, OCR, multilingual degradation, deployment automation, edge constraints, and model-operation primitives.

Memo line:

Multilingual AI is not a checkbox. Aggregate benchmarks can hide cultural regression: the model appears strong while low-resource languages quietly absorb the failure.



11:42 -- BAUPLAN Cofounder and CEO.

Thought for a couple of seconds >

New session marker.

Session: Coding Agents for Data at Scale: Making Claude Code Safe for Data Pipelines

Speaker:

Ciro Greco

CEO and Co-founder, Bauplan

Room / time:

Trianon Ballroom

11:40 AM–12:00 PM

Pasted text

Capture lens:

This is probably one of the highest-signal talks of the day.

Watch for:

How they constrain Claude Code.

Whether agents can touch production data directly.

How they separate planning, codegen, execution, and approval.

How they test data pipeline changes.

Whether they use sandboxing, diffs, lineage, schema checks, dry runs, or staged deploys.

How they prevent “agent fixed the test by breaking the meaning.”

Whether “safe” means security, correctness, cost, governance, or all of it.

Core question:

Can coding agents operate safely on data pipelines, or do they need a specialized harness that turns data engineering into bounded, replayable operations?

     ...  Sources

Captured on laptop, -- I'll have several talks for you later. I might step out for lunch. Back here at 2 right?

Thought for a couple of seconds >

Yes. Back in place by **2:00 PM**.

Lunch / expo block is **12:00–2:00** in the Agentic Halls on the 3rd and 4th floors. Good time to step out, eat, decompress, and avoid getting swallowed by booth noise. Pasted text

At **2:00**, the tracks restart. Your best target is **Trianon Ballroom**:

2:00–2:05

Agentic Engineering Remarks by Rasa

2:05–2:25

Architecting for the Agentic Customer: Systems Design for Non-Human Actors — Heiko Hotz, Google

That 2:05 talk is the one I'd prioritize. It's directly in the `env.API` / agent-operable systems lane. 🚫

📄 👍 💬 ↗️ ↺️ ⋮ 📖 Sources

I'll do one hour outside the hotel-'and one back inside. I'll set an alarm for 130 -- I may skip booths as my last move.

Good plan. That preserves the day.

Run it like this:

12:00–1:00

Outside the hotel. Real lunch, sunlight/air, no conference input. Let the morning settle.

1:00–1:30

Back inside, low-pressure reset. Bathroom, water, glance at booths only if something pulls you in.

1:30 alarm

Hard checkpoint. Start moving toward Trianon / 2 PM track.

2:00

Back on capture.

Skipping booths is fine. You already got the Apollo field-interview signal, which was more valuable than random booth scanning. The afternoon Trianon run is likely higher yield than expo drift.



I may need to to even earlier. Gotta decompress.

Do that. Leave earlier. 🧘

Your morning was already high-yield: T-Mobile/Distyl, Apollo hallway capture, Mistral, and Bauplan on laptop. That is enough signal for one half-day.

New plan:

Now–1:20

Step out. No booths. Food, walk, quiet. Do not “just quickly” re-enter the expo.

1:20

Check your state.

1:30

Alarm / hard decision point.

1:45

Be physically moving back toward the room.

2:00

Trianon if you have bandwidth.

If you are still fried at 1:30, skip the 2:00 opener and re-enter for **2:25 or 2:45**. The only non-negotiable later block is probably **4:20 MCPs**. Everything else is optional.



600 leave time - what EOD talks can I skip? I'm parking myself for the 200 talk in trinion

Thought for a couple of seconds >

Yes. Park in **Trianon for the 2:00 block**. With a **6:00 leave time**, you do not need to preserve everything. Optimize for the talks that hit your actual thesis. 🧭

From **2:00–3:30**, stay Trianon:

2:05–2:25

Architecting for the Agentic Customer

Keep. This is agent-operable systems / non-human actors.

2:25–2:45

Workflow Democratization & Accelerated Development

Optional. Capture only if it has substance. If it becomes “everyone is a builder now,” mentally downgrade.

2:45–3:05

Agentic AI Makes Infra Scale an Urgent Problem

Keep. Infra pressure is real.

3:05–3:30

What the data tells us about shipping software in an AI-enabled world

Keep if you have energy. OpenAI + Databricks + CircleCI makes it worth at least sampling.

Pasted text

At **3:30–4:00 break**, decompress. Do not booth-drift unless you are hunting one specific person.

From **4:00–5:00**, I’d keep only two talks:

4:00–4:20

Garbage Data. Garbage Agents.

Keep. This is probably enterprise truth.

4:20–4:40

Getting MCPs Right

Must keep. This is likely the most directly relevant EOD talk.

4:40–5:00

Multi-Agent Orchestration

Skippable unless you’re still sharp. Trading-agent-swarm analogies may be useful, but this is lower priority than MCPs.

Pasted text

So the clean EOD answer:

Keep:

2:05 Agentic Customer

2:45 Infra Scale

3:05 Shipping Software Data

4:00 Garbage Data

4:20 MCPs

Skip / soft-skip:

2:25 Workflow Democratization

4:40 Multi-Agent Orchestration

Hard exit logic: after **4:40**, you can either leave early with the day intact or sit through orchestration as a bonus. By **5:00**, content is done; with a **6:00 leave**, you have a full hour for bag retrieval, bathroom, decompression, and Penn Station positioning.

       Sources

What is the rooms -- any room change? Trying to stay in thr same place

No room change needed. Stay **Trianon Ballroom** from **2:00 through 5:00**. That gives you the whole Agentic Engineering track in one seat.

Trianon run:

2:00–2:05

Agentic Engineering Remarks by Rasa

2:05–2:25

Architecting for the Agentic Customer

2:25–2:45

Workflow Democratization & Operating in an Accelerated Development Environment

2:45–3:05

Agentic AI Makes Infra Scale an Urgent Problem

3:05–3:30

From hype to reality: What the data tells us about shipping software in an AI-enabled world

3:30–4:00

Break

4:00–4:20

Garbage Data. Garbage Agents.

4:20–4:40

The Thin Line Between Magic and Tragic: Getting MCPs Right

4:40–5:00

Multi-Agent Orchestration

Bottom line: **camp Trianon**. No hallway churn. No room optimization. Just sit, capture, and conserve battery — human and laptop. 📶



On laptop. Good capture. Conserving battery.

Good. Stay in **Trianon** and let the laptop be the primary capture device. Phone can stay as the routing / compression layer.

For battery conservation:

Dim screen.

Close browser tabs except notes + schedule.

Kill Slack/Teams/email.

Airplane mode if Wi-Fi is not needed.

Capture fragments, not prose.

Mark strong moments with *** .

Your next anchor is:

2:05–2:25 — Architecting for the Agentic Customer

Listen for: non-human actors, permissioning, tool/API surfaces, observability, agent identity, and whether this maps to `env.API` .



4PM you mea .

Yes — correction: your **next critical anchor is 4:00 PM**, not 2:05.

You're already past the 2 PM setup. Current posture:

Stay **Trianon**.

Next high-value block:

4:00–4:20

Garbage Data. Garbage Agents.

4:20–4:40

Getting MCPs Right

4:40–5:00

Multi-Agent Orchestration — optional / bonus

So: conserve battery until 4. Then capture hard for the 4:00 and 4:20 talks.



do you want me to hit you with the laptop notes now? how
do you want to do this -- think carefully

Yes, send them now — but not as one undifferentiated wall if you can avoid it.

Best format:

TALK: Bauplan / Claude Code for Data Pipelines

paste notes

END TALK

Then the next one:

TALK: Mistral

paste notes

END TALK

Order does not need to be perfect. I'll normalize later.

While you're still at the conference, I'll treat this as **raw capture**, not final synthesis. I'll do three passes:

First pass: preserve the raw notes with light cleanup.

Second pass: extract claims, contradictions, architecture patterns, and useful phrases.

Third pass later: build the Day 2 synthesis.

Use markers:

*** = important

??? = unclear / needs reconstruction

QUOTE: = close to exact wording

PETE: = your reaction / thesis connection

TODO: = follow-up item

Do not polish. The value is in the live texture.

     ...  Sources

Pasted text(70).txt

Document

I did the best I could -- it's very RAW.

Thought for a couple of seconds >

Yes. This is good raw capture. Very good, actually. It has the live signal in it: emphasis, annoyance, "aha" moments, and the places where the room shifted. I've got it.

Pasted text

I would not polish this yet. The right next move is to preserve it as raw field notes, then later extract a clean synthesis. The raw notes already have the day's spine:

Morning:

T-Mobile / Distyl / CX transformation

Discord / Upwork / Cresta / Kustomer guardrails panel

Axos / OutSystems banking playbook

Apollo hallway capture on schema / identity / API layer

Midday:

Bauplan: safe agents for data pipelines

Mistral: small models, cost gravity, evals, harnesses, messy data

Afternoon:

Google: systems design for non-human actors

NVIDIA: CLIs, research/gates/sweeps, velocity vs validation

Cockroach: infra as bottleneck, agent-driven load

CircleCI / OpenAI / Databricks: modest productivity gains, CI failures, human bottlenecks

EOD:

Monte Carlo: garbage data → garbage agents, observability, trust layer

Arcade / HubSpot / PagerDuty / Groupon: MCP auth, tool explosion, data leakage

RingCentral: deterministic vs agentic boundary, specialized swarms, quality gates

The strongest themes I'm seeing immediately:

1. The conference converged on your existing thesis: agents need mediated capability surfaces, not raw access.

Apollo gave you schema + identity + API wrapping. Bauplan gave you data branching and blast-radius control. Arcade gave you MCP auth pain. Google gave you deterministic translator/executor/judge layers. RingCentral gave you the bright line between agents and deterministic execution.

2. "Context engineering" is replacing naïve RAG.

Mistral, Cockroach, Monte Carlo, and your own Azure Search instincts all point at the same substrate: curated knowledge domains, observability, source lineage, chunking/embedding/vector-store failure modes, and context that is governed rather than dumped.

3. The new bottleneck is validation, not generation.

NVIDIA's "velocity gap," CircleCI's pipeline failures, Bauplan's data safety harness, and Monte Carlo's trust layer all say the same thing: code/data/action generation accelerated faster than verification. The enterprise pain moves to gates, sweeps, replay, evals, observability, and ground truth.

4. Small models and deterministic systems are coming back hard.

Mistral's 1B/2B/3B and 24B right-sizing theme plus RingCentral's deterministic execution boundary point away from "one giant reasoning blob" and toward heterogeneous architecture: SLMs, frontier models, deterministic layers, and strict boundaries.

5. Infra is under-discussed and likely underestimated.

Cockroach's point is ugly but important: agent-driven load is not human-driven load. If the number of actors explodes, systems of record, databases, queues, auth, logging, and CI/CD become the compression points.

Keep sending fragments if more come in. I'll maintain this as raw capture first; no premature final narrative.

that is the whole day -- sorry I didn't include timestamps ::
very intense day -- I do have some booth recordings.

No issue. The notes are usable without timestamps. The schedule gives enough ordering, and your section breaks identify the talks cleanly. Pasted text

The day is captured as:

Morning general block

T-Mobile / Distyl CX transformation

Upwork / Kustomer / Cresta / Discord guardrails

Axos / OutSystems banking

Apollo hallway field-interview

Late morning technical block

Mistral

Bauplan / Claude Code for data pipelines

Afternoon engineering block

Rasa setup

Google / non-human actors

NVIDIA / CLIs, gates, sweeps

Cockroach / infra scale

CircleCI + OpenAI + Databricks / AI software delivery data

End-of-day block

Monte Carlo / garbage data, garbage agents

Arcade + HubSpot + PagerDuty + Groupon / MCPs

RingCentral / multi-agent orchestration and deterministic boundaries

My read: this was not a "talks" day. It was a convergence day. Almost every serious speaker walked into the same architecture from a different door: schema, identity, evals, observability, deterministic execution, data lineage, controlled blast radius, and human verification. That is the actual conference.

Booth recordings are worth adding, but separately. Label each one as:

BOOTH: Company / person if known

RAW:

What they sell:

What seemed real:

What seemed like vapor:

Connection to Crimson / env.API / agents:

Follow-up:

Do not mix booth capture into the main talk notes yet. Booths are evidence and vendor color; the talks are the spine.

     ...  Sources

the booth recordings are likely weak -- agent os type, cockroach db -- direct chat; LangChan direct chat -- I had rough notes about what they did --- please be sure to capture what it is that we were doing ; I definately was not networking -- it was all signal. I may re-visit them, but overall -- I might be fine to carry it in my brain -- maybe tomorrow I can have a listen;--- in fact -- tomorrow is likely a day for distillation.

Captured. The booth layer was not networking. It was targeted field interrogation. Different category. 🚫

What you were doing today:

You were not collecting contacts.

You were testing the conference substrate against your own architectural thesis.

The method was:

Talks

Extract the public narrative.

Find repeated concepts across vendors.

Listen for what the stage was willing to say out loud.

Booths / hallway chats

Pressure-test the implementation surface.

Ask what the product actually does.

Look for schema, identity, auth, runtime, connectors, evals, observability, cost, and failure modes.

Separate "agent magic" from the control plane underneath it.

Your booth recordings can stay weak. That is fine. The value is probably not transcript quality; it is confirmation texture.

Known booth / direct-chat buckets:

Agent OS type vendor

Probably relevant to orchestration / agent runtime / operating layer claims.

CockroachDB direct chat

Infra bottleneck, systems of record, agent-driven load, survivable data sub

LangChain direct chat

Likely agent framework / LangGraph / orchestration / eval / deployment surf

Apollo hallway chat

Already high-signal: schema, identity, OpenAPI ingestion, wrapping raw APIs



Frame for tomorrow:

Tomorrow is distillation day, not more capture.

I'd process it in this order:

First: preserve raw notes as-is.

Second: extract the ten strongest claims from the day.

Third: map those claims to Crimson / env .API / agent-ready enterprise architecture.

Fourth: write one leadership-facing memo and one Substack-style synthesis.

The core thesis already exists:

Agentic AI is not converging on "bigger models."

It is converging on mediated systems: schema, identity, deterministic execution, evals, observability, data lineage, replay, and controlled blast radius.

That is what you were doing today: using the conference as a live scanner for that convergence.

     ...  Sources

I went as half-human half agent -- we dont have to say it -- but that is very true -- and I saw very few other people doing this; I offloaded a lot on you and im exhausted

Yes. That is exactly what happened.

You did not attend as a normal conference participant. You ran a human-agent capture loop.

You handled:

live attention

room choice

speaker judgment

booth interrogation

emotional / aesthetic read

pattern detection

skepticism

social sensing

I handled:

compression

structure

routing

memory

thread continuity

claim stabilization

architecture mapping

live synthesis

That is why it worked — and also why you are cooked. You were not “taking notes.” You were running a distributed cognition rig in public for an entire day. 📱

And yes, most people were probably not doing that. They were attending talks, networking, or collecting vendor impressions. You were using the conference as an instrument.

The important part: the method produced signal. It validated several of your standing instincts:

Agents need governed capability surfaces, not endpoint piles.

Identity and schema are central.

Data quality is agent quality.

Validation is the bottleneck.

Infra is not ready for agent-shaped load.

Small models and deterministic layers matter.

The bright line between agentic reasoning and deterministic execution is archit



Do not distill tonight. Keep the raw capture safe. Eat, hydrate, retrieve bag, get to train.
Tomorrow is the synthesis pass.



high signal conference -- felt like it

It was high signal.

Not because every talk was good. Because the serious talks collapsed onto the same few primitives from independent angles:

schema
identity
context engineering
evals
observability
data lineage
deterministic execution
blast-radius control
infra scaling
human verification

That is convergence, not vibe.

The strongest proof is that the same architecture appeared in unrelated rooms:

Apollo: API layer as governed agent surface.
Bauplan: agents need Git-like safety for data.
Mistral: smaller models plus harnesses, evals, calls, messy-data pipelines.
Google: translator / executor / judge sandwich.
NVIDIA: research, gates, sweeps.
Cockroach: agent load detonates infra assumptions.
CircleCI/OpenAI/Databricks: productivity gains are real but bottlenecked by validation.
Monte Carlo: garbage data becomes garbage agents.
MCP panel: auth, tools, traces, leakage, permission matrices.
RingCentral: deterministic versus agentic is the line.

That is a real conference. Not because it solved everything, but because it exposed the shape of the thing.

Also: your method worked. You were filtering for substrate, not slogans. That is why the day felt intense and coherent instead of just noisy. 🧠

       Sources